

Reproduction and Analysis of Sparse Online Learning via Truncated Gradient

Siddhartha Shukla

Roll Number: 230104

Rishihood University

Advanced Machine Learning Mid-Semester Examination

Abstract—This report presents a reproduction and analysis of the paper “Sparse Online Learning via Truncated Gradient” by Langford, Li, and Zhang. The paper introduces the Truncated Gradient algorithm for estimating sparse online models. We reproduce the key contribution using a synthetic dataset, analyze the results, conduct ablation studies to understand the role of different components of the method, and investigate failure cases when the underlying assumptions of the model are violated.

I. INTRODUCTION

Understanding relationships and maintaining efficiency in high-dimensional real-time data is an important problem in machine learning. Online gradient descent models typically suffer from dense weights that consume excessive memory. The paper by Langford, Li, and Zhang proposes the Truncated Gradient, an algorithmic method that successfully induces feature sparsity by actively shrinking and truncating small weights.

II. PAPER SUMMARY

The Truncated Gradient algorithm solves the lack of actual sparsity mapping in L1-regularized stochastic gradient descent (SGD). Standard stochastic gradient updates consist of non-zero floating-point operations that mathematically rarely sum to precisely zero.

To resolve this, the algorithm actively shrinks small weights towards zero at a defined gravity rate (g) and truncates them completely if they fall below a threshold (θ). This allows models to learn only the truly informative features, yielding massive computational and storage savings with rigorously proven, bounded minimal regret to prediction accuracy compared to generic SGD.

III. REPRODUCTION SETUP

To reproduce the core idea of the Truncated Gradient algorithm, a synthetic dataset was generated with high inherent sparsity. The dataset contains:

- 1,000 samples
- 100 features
- 10 mathematically informative features (90% true sparsity)

The experiment was implemented in Python. We developed the Truncated Gradient logic exactly as detailed in Algorithm 1 for Least Squares from the original text. For comparison, we also implemented a standard SGD baseline.

IV. EXPERIMENTAL RESULTS

The resulting weight vectors showed significant sparsity properties, indicating that the algorithm successfully identified irrelevant features and truncated their influence to zero.

- The SGD baseline achieved 0.0% true sparsity, memorizing noise distributions alongside the valid signal with a test MSE of ≈ 0.027 .
- The Truncated Gradient reproduction successfully drove the 90 irrelevant feature weights identically to 0.0, matching the 90% true sparsity metric while maintaining comparative predictive MSE.

This result strongly substantiates the core claim and mechanics reported in the original paper.

V. ABLATION STUDY

To understand the contribution of different components of the method, two ablation experiments were performed on the model.

A. Removing the Threshold

In the first ablation experiment, the threshold component was removed by setting $\theta = \infty$. This transforms the algorithm into a standard naively L1-regularized online gradient descent. By subjecting all weights—even highly informative ones—to the gravity shrinkage, the model critically under-fits the strong signals, degrading the predictive representation.

B. Removing Gravity

In the second ablation experiment, the gravity parameter was eliminated by setting $g = 0$. This effectively removes the core truncation mechanic entirely. As expected, the resulting sparsity collapses exactly to 0.0%, confirming that the active gravity pull is the sole enabler of the algorithm’s functional sparsity design.

VI. FAILURE MODE ANALYSIS

The Truncated Gradient algorithm fundamentally assumes that the online dataset features have substantial inherent sparsity. To test a failure scenario, a completely dense dataset was constructed where all 100 features were robustly informative and contained zero noise variables.

When the algorithm was applied to this pure, dense signal, the truncation penalty actively destroyed valid mapping relationships. The enforced constraint of driving small weights

to zero forced the model to under-fit, drastically worsening the MSE compared to a generic SGD baseline. This explicitly proves that when the underlying assumption of sparsity is actively violated, the forced-sparsity algorithm is mathematically detrimental.

VII. REFLECTION

The reproduction exercise provided clear insight into how continuous algorithms can enforce rigid structural behavior on continuous variables during massive online stream processing.

The most surprising aspect of this implementation was the sheer architectural simplicity of the core truncation logic—a basic `max()` and `min()` step conditional—compared to the massive formal algebraic proofs required natively to guarantee its stable regret error bounds. If time permitted, reconstructing the dynamic hash-table “lazy-update” mechanic outlined in Section 5 of the paper would be the optimal next step, mathematically fulfilling the operational metric of true $O(k)$ time complexity processing.